RTMSF: Event-Driven Partition Membership Maintenance for Existing Cloud Data Lakes

Ankit Joshi · Independent Researcher, USA

October 2, 2026

Abstract

Independent ingestion, backfill, and cleanup jobs can change ordinary object-store files without updating the partition membership of an existing Hive-style catalog. Scan-based repair discovers missing partitions, while event-based crawlers can limit discovery to affected folders. This paper presents the Real-Time Metadata Synchronization Framework (RTMSF), a provider-neutral system specialized for known-schema partition membership. WOERA groups notifications by partition using processing-time settling windows; PMCP checks current storage and performs idempotent catalog additions or removals. Unacknowledged notifications supply replay after worker failure. An AWS implementation is evaluated against Athena MSCK and a Glue event-based crawler using valid Parquet data, independent catalog observations, and corroborated S3 request measurements. Across the tested 10–1,000-partition range, RTMSF used one HEAD request for an isolated new partition. For an append to a registered partition containing 100 files, RTMSF used one LIST request, while the configured crawler used 113 HEAD/GET/LIST requests. A controlled per-event handler comparison required 20 storage checks for 20 quick uploads to one partition, versus one for RTMSF; both required 20 when uploads targeted 20 different partitions. Active-load trials retained progress and correct final state, but a single worker accumulated backlog at 20 changes/s across 1,000 partitions. Increasing its buffer improved catch-up without establishing sustained capacity at that rate. Athena and the immediate handler exposed new partitions sooner in their tested comparisons. The evidence supports bounded storage inspection and workload-dependent coalescing, with an explicit delay and capacity tradeoff, rather than universal performance or cost superiority.

Keywords: partition membership; metadata reconciliation; object storage; Hive catalog; event coalescing; crash recovery.

1. Introduction

Hive-style external tables describe data stored beneath partitioned object-store prefixes. A query engine can read files in a registered partition, but a newly created partition may remain invisible until its catalog entry is added. Conversely, deleting every file in a partition can leave a stale catalog entry. Storage and catalog changes are separate operations when ingestion and cleanup jobs write ordinary files without participating in a shared table commit protocol.

This paper addresses that deployment condition. Table identities, schemas, and partition layouts are already known. Writers need not change their data format or call the catalog. The task is to maintain membership: register nonempty partitions and remove registrations for empty ones. It does not include discovering new tables, inferring schemas, providing table transactions, or indexing records.

Consider a catalog-driven sales-data dispatcher. Independent ingestion jobs write files into tenant/hour partitions, while retention jobs delete expired data.

The dispatcher enumerates the catalog to discover partitions and schedule work. A missing registration delays discovery of new input; a stale registration can schedule a task that finds no files. Membership maintenance therefore serves the dispatcher's existing discovery interface. This is a motivating workflow, not a deployed consumer measured here. Membership identifies nonempty storage; batch readiness and the presence of queryable rows remain separate concerns.

The intended deployment needs enumerated membership in an existing catalog, potentially shared by multiple consumers. An Athena-only deployment with predictable partition values may avoid this task through partition projection [12]. Empty registrations ordinarily return no rows; their removal serves accurate catalog enumeration and metadata maintenance. We do not measure a query-result error caused solely by retaining an empty registration, or quantify production demand for removal. Multi-engine consumption is a deployment motivation, not a measured result of the present AWS query checks.

Athena MSCK discovers and registers missing Hive-style partitions by inspecting storage [1]. Glue's event-based crawler consumes notifications to restrict subsequent crawls to affected folders after initialization [2]. Both are relevant alternatives. An event-based crawler already provides work locality, so comparing only against a full scan would miss the stronger question: how much inspection is needed when the schema is known and only partition membership must be decided?

RTMSF uses an object notification to identify a partition worth checking. The notification does not establish its current state. One surviving object establishes nonemptiness; an empty-prefix observation justifies attempting removal within the system's stated concurrency boundary. This permits bounded storage checks without reading the files to discover their contents. Coalescing reduces repeated checks when many notifications concern the same partition.

The contribution is the complete membership-maintenance system and the evidence for its design choices:

- An integrated design for existing catalogs and independently managed writers, combining partition routing, event coalescing, bounded storage verification, catalog mutation, and delivery acknowledgement through provider-neutral interfaces.
- An explicit processing and recovery contract for the combined WOERA and PMCP system, including retained-message replay, bounded retries, and limits imposed by non-atomic storage and catalog operations.
- An AWS evaluation with retained reproduction inputs that separates namespace traversal, affected-folder file inspection, and repeated checks of the same partition. Two deployed AWS alternatives and a controlled per-event comparison establish the measured work/delay tradeoffs; lifecycle, crash, active-load, and buffer checks exercise the system's behavior and operating limits.

WOERA and PMCP describe complementary parts of this system. Event grouping, existence checks, retries, idempotence, and replay are established techniques. The proposed contribution is their integration for this external-catalog task, the explicit processing contract, and the measured consequences of specializing maintenance to membership. Portability is a design property; the evaluated implementation is AWS only.

2. Background and related work

Discovery and repair. Athena's MSCK REPAIR TABLE discovers missing partitions and adds them to the catalog; it does not remove stale partitions [1]. Its request count depends on the namespace it traverses. A partition count is therefore not an implementation-independent prediction of one LIST per partition.

Glue event-based crawling uses S3 notifications delivered through SQS after an initial crawl. Events identify affected folders, but a subsequent crawl must

still be started by API or scheduling [2]. Our comparison uses an initialized event crawler with an existing catalog-table target and schema-preserving settings. It measures that configuration, not every Glue mode.

Ordinary-file automatic refresh. Snowflake external tables can refresh file metadata from S3 notifications through SQS, adding new files, updating changed files, and removing files no longer present [10]. External-table partition expressions can derive partition values from paths [11]. This is directly relevant precedent for metadata refresh without a writer-managed table commit protocol. It maintains Snowflake external-table file/partition metadata rather than providing an adapter for an existing shared Hive catalog. Its documentation does not establish RTMSF's bounded-check behavior; this paper provides no performance comparison with it.

Avoiding registration. Athena partition projection derives partition values from configured rules and ignores registered partition metadata [12]. For suitable Athena-only workloads, including predictable layouts like our fixtures, this can remove the need for registration maintenance. It does not maintain an exact partition set for other consumers of the shared catalog. We therefore distinguish this deployment alternative from the measured registration baselines.

Format-managed metadata. Delta Lake maintains table metadata through its transaction log and commit protocol [3]. Hudi's metadata-as-data design supplies metadata and indexes within Hudi-managed tables [4]. These systems offer broader table-management capabilities and are architectural alternatives when writers can adopt their protocols. Their published query/index benchmarks do not measure the same task as reconciling an existing catalog after arbitrary ordinary-file writes. Their different approach is relevant to deployment choice, rather than a reason to classify all metadata management as the same problem.

Catalog infrastructure and recovery mechanisms. TreeCat addresses catalog storage, querying, and transactions; its discussion identifies integration of physical data and metadata consistency as a remaining challenge [5]. RTMSF operates against an existing catalog's partition APIs. MillWheel provides relevant precedent for durable processing and failure recovery [6]. RTMSF's external effects remain idempotent operations with replay, rather than exactly-once execution. These systems inform the design but are not interchangeable measured baselines.

Approach	Where metadata maintenance originates	Relationship to this work

Athena MSCK	On-demand or scheduled discovery of storage paths	Measured ADD repair baseline; no stale-partition removal
Glue event-based crawler	Notifications select folders for a started crawl	Closest measured event-based alternative for ordinary files
Snowflake external-table refresh	Notifications update external-file metadata and derived partitions	Related automatic-refresh service for Snowflake's catalog
Athena partition projection	Query-time derivation of configured partition values	Can avoid registration for suitable Athena queries; does not maintain shared catalog membership
Delta Lake / Hudi	Table-format commit and metadata protocols	Architectural alternatives requiring compatible writers
TreeCat	Catalog operations and internal metadata organization	Complementary catalog infrastructure
RTMSF	Notifications, partition-level coalescing, and bounded current-storage checks	Known-schema membership maintenance with retained-input replay

An event handler is the code that receives a notification and processes the resulting work. An immediate handler checks each event separately; a grouped handler combines events for the same partition before checking. Section 5.7 compares RTMSF with an immediate control using the same worker, PMCP, adapters, retries, and acknowledgement rules. This isolates the effect of grouping within the implementation; the control is not an independently optimized competing system. A straightforward grouped handler with equivalent timing, bounded storage checks, ownership, retry, and delivery-completion rules could provide the same central storage-work bound and retained-input recovery behavior. RTMSF specifies and implements that general design; no fundamental advantage over an equally capable grouped implementation is established here. This is an analytical comparison to a design pattern, not another measured or published system. The systems claim rests on the integrated design and its evaluated tradeoffs. The literature comparison establishes scope and responsibilities, rather than proving that no earlier system could implement this contract.

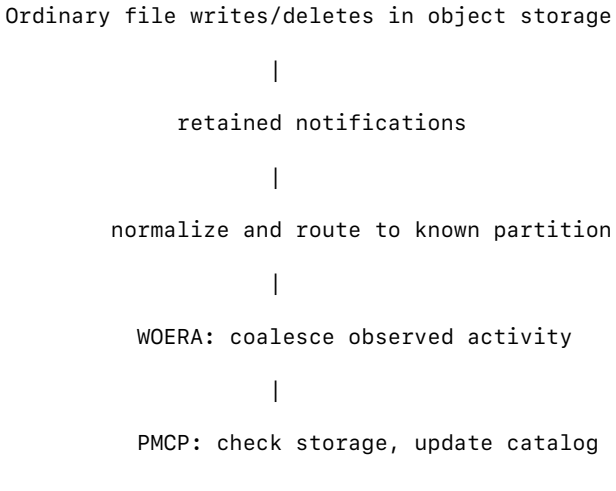
3. System model and design

3.1 Membership and assumptions

For a configured partition prefix p, let S(p) indicate that storage contains at least one object under that prefix and C(p) indicate a catalog registration. The reconciliation target after relevant activity settles is C(p) = S(p).

This is a target state, not a transactional invariant during concurrent writes.
Initial alignment is assumed or must be supplied by a separate repair procedure;
an untouched drifting partition is not visited by this worker. Convergence requires
an adequate eligible notification after the relevant final storage-state change,
followed by successful processing within the retention/retry limits.
The current implementation routes CREATE and DELETE notifications for every
object beneath a fully matched configured partition path, including markers and
nested objects. Partition values and decoded object keys are preserved exactly.
Its emptiness check counts any remaining object. A marker-only prefix therefore
counts as nonempty; membership does not assert that queryable rows exist.
The source used for the five primary figures routed only Parquet-file events.
That version could leave stale membership after a final _SUCCESS deletion.
Routing was subsequently aligned with PMCP's all-object existence check, without
changing WOERA or PMCP. The separate validation run exercises marker-only
registration, deletion of Parquet while a marker remains, and removal after the
marker is deleted (Section 5.6). Primary figures use marker-free fixtures;
their older source is retained and is not presented as covering this correction.
The contract assumes that relevant notifications are admitted to the configured
event source and retained long enough for processing, routing is valid, and the
worker and providers are available for retries to make progress. The AWS deployment
uses one active worker. Permanent notification omission, queue purge or expiry,
permanent provider failures, and external schema changes are outside the contract.
Checks of storage and changes to the catalog are not atomic with arbitrary writers.

3.2 Architecture and routing



```
                |
complete all work or durably record terminal input
                |
            acknowledge
```

The core uses four interfaces: event source, object store, partition catalog, and telemetry sink. At startup, the AWS catalog adapter resolves configured table locations and partition columns. Routing rejects overlapping roots and maps each eligible object key to one table and partition. Configuration changes require a controlled restart. The current routing implementation iterates over configured tables; it is not a demonstrated constant-time lookup over a large table fleet.
The public data flow is ObjectEvent → PartitionEvent → PartitionIntent.
Delivery handles remain separate from these algorithmic values. Source timestamps remain optional event metadata; the worker's monotonic clock determines deadlines.
The AWS adapters use S3, Standard SQS, and Glue; an ECS/Fargate task runs the worker.
The worker depends on the EventSource interface, whose operations receive, acknowledge, renew, and durably record terminal deliveries. AWS receipt handles and SDK calls remain inside SqsEventSource; aws_app.py supplies the wiring.
Acknowledgement accepts a sequence of deliveries, without requiring a native batch API. An alternative adapter may implement it with individual calls, within its declared call/lease budgets and failure semantics. This separation supports porting; it does not establish conformance of an unimplemented cloud adapter.

3.3 WOERA: processing-time coalescing

WOERA, the Weakly Ordered Event Reconciliation Algorithm, maintains one active record per observed partition. A record holds its first observation time, current deadline, candidate mutation, observation count, and at most three observed CREATE keys. For an arrival at local time t, initial eligibility is:

```
eligible = min(t + Q, first_observation + T_max)
```

The quiet interval Q restarts on an observed arrival. T_max prevents continuous arrivals from indefinitely extending initial buffering. A scheduled retry imposes a separate lower bound on eligibility; a new arrival cannot bypass that retry.
These are processing-time rules, not event-time watermarks or a bound on end-to-end

visibility. Worker scheduling, delivery, provider calls, and retries add time.
Within a window, any CREATE makes the candidate ADD; otherwise the candidate is DROP. This ADD-dominance rule aggregates observations without assuming their source order. PMCP still consults storage, so the candidate is not a command to blindly create or delete a registration. After completion, later notifications can open another window for the same partition.
A versioned deadline heap supports scheduling; stale entries are periodically compacted. The worker limits pending deliveries, events, and partitions to avoid unbounded intake. Bounded coalescing state does not mean the system can accept an unlimited arrival rate. The AWS adapter additionally limits retained raw message bytes and stops fetching when capacity is unavailable; excess input remains in SQS. The default limits are 100 deliveries, 1,000 active partitions, 10,000 decoded events, and 16 MiB of raw payload. Parsed events, table definitions, and SDK memory require additional RAM. Sections 5.8–5.9 report bounded activity and overload observations; retry storms and many-table throughput remain unmeasured.

3.4 PMCP: membership decisions

PMCP, the Partition Mutation Commit Protocol, uses current storage observations and catalog state. The following table describes its decision paths. A prefix check requests at most one listed object; file contents are not inspected.

Candidate and catalog state	Storage observation	Action
ADD, absent	An observed CREATE object exists by HEAD, or fallback prefix is nonempty	Create registration
ADD, absent	No observed object or prefix content found	Retry; eventual exhaustion is terminal
ADD, registered	Prefix nonempty	Keep registration
ADD, registered	Prefix empty	Remove stale registration, then retry
DROP, registered	Prefix nonempty / empty	Keep / remove registration
DROP, absent	Prefix nonempty / empty	Create registration / no change

The absent-ADD path checks at most three observed keys before one bounded LIST fallback. The measured normal new-partition path succeeds with its first HEAD. AlreadyExists and NotFound responses are treated idempotently. Transient provider errors become retries; fatal errors receive terminal handling. The word “commit” in the protocol name denotes completion of this processing step, not an atomic

transaction spanning object storage and the catalog.

The storage-work argument is local: existence can be established by one object, whereas absence requires checking the prefix. For A PMCP attempts, bounded HEAD and prefix checks give O(A) storage-check invocations, independent of unrelated partition count. A includes retries and later windows for the same partition; it is not always the number of distinct changed partitions. Every event must still be consumed, and provider-side LIST cost is not inferred from this API-call bound. Catalog, queue, compute, and telemetry work are additional costs.

3.5 Delivery completion and crash recovery

The worker retains each delivery until every partition intent derived from it has completed or the terminal input has been durably copied to the dead-letter queue. It then acknowledges the delivery. Leases are renewed while work is pending. An unsafe lease or detected operation-budget overrun terminates the worker for replay; SDK timeouts are not a hard guarantee against every OS or network stall.

A crash discards in-memory coalescing state. Unacknowledged input reappears from SQS and opens fresh windows. If a catalog call succeeded before the crash, replay may attempt it again; current-state checks and idempotent catalog responses handle that uncertainty. No separate checkpoint database is required by this deployment.

This is not exactly-once mutation execution or a distributed fencing protocol. After storage has stopped changing, a successful appropriate observation and catalog update can establish the desired membership. During concurrent activity, a writer may change storage between a check and its catalog operation. Subsequent admitted notifications can trigger reconciliation, subject to retention, retries, and terminal failure handling. A stale CREATE whose file has disappeared can exhaust retries and enter the dead-letter queue even when the catalog is correctly empty. The system must not be described as error-free for every event ordering.

4. Experimental methodology

4.1 Implementation and fixtures

All five figures use execution run_20260930_004034_214f, initiated September 30, 2026 UTC in AWS us-west-2. The worker uses one ARM64 Fargate task with 0.25 vCPU and 512 MiB, Q = 10 s, T_max = 31 s, a 180 s SQS visibility lease, and at most

five retries after an initial PMCP attempt. Source and image identities, trial settings, service responses, query results, and artifact hashes are retained. All recorded MSCK queries selected AUTO and used effective Athena engine version 3.

The main layout is:

```
orders/dt=2026-01-01/site=site-000/hour=00/part-000001.parquet
```

All partition columns are strings. Hours 00–23 are filled before moving to the next site, and dates advance after 100 sites. Each file is valid Snappy Parquet with one row: record_id and amount are BIGINT, with amount = 10 × record_id. Corresponding method fixtures use identical file bytes. Every comparison condition, method, and repetition receives a fresh table and bucket. Existing partitions are registered before measurement; initialization, setup notifications, and baseline query checks finish before the measured change.

The evaluated matrix is:

Output	Conditions actually run	Repetitions	Comparison trials
Figures 2 and 3, shared trials	10, 100, 500, 1,000 starting partitions; three methods	2 per condition/method	24
Figure 4	1 or 100 existing files in one partition; 100 starting partitions; two methods	2 per condition/method	8
Figure 5	10, 100, 500, 1,000 starting partitions; two layouts; two methods	2 per condition/layout/ method	32
Figure 1	First file, append, partial deletion, final deletion	1 lifecycle	—
Recovery and supporting checks	Four crash cases and seven check types	1 per case/type	—

There are 64 distinct comparison trials. Figures 2 and 3 reuse the same 24 trials; they are not independent experiment sets. Method order rotates between repetitions, and layout order reverses. Two repetitions describe observed variation but do not support tail-latency estimates, confidence intervals, or stable service rankings.

4.2 Baselines and measurement endpoints

RTMSF runs continuously. Athena receives an actual MSCK REPAIR TABLE command promptly after the new file upload. The Glue crawler targets an existing catalog table in CRAWL_EVENT_MODE, preserving table schema and inheriting it for partitions.

It completes its initial crawl and starts a subsequent crawl once the new notification is available. Startup and initialization work are excluded from steady-state request measurements for all methods.

The crawler has UpdateBehavior=LOG, DeleteBehavior=LOG, and partition AddOrUpdateBehavior=InheritFromTable. Its deletion policy logs changes rather than removing stale registrations. The cross-method trials therefore compare ADD and append work; full ADD/DROP lifecycle correctness is tested only for RTMSF.

Figure 3 ends at the first successful independent Glue GetPartition observation, polled approximately once per second. Its method-specific starts are:

Method	Timer start
RTMSF	Successful upload response
Athena MSCK	Immediately before StartQueryExecution
Glue event crawler	Immediately before StartCrawler, after notification availability

Submission latency, queueing, and execution after these starts remain included. RTMSF's timing additionally includes event delivery and its intentional settling window. The graph does not measure a common upload-to-visibility interval or a production scheduling delay. Polling measures first observed visibility, not the exact catalog-write instant. An independent SELECT verifies query rows afterward; SELECT completion is not the timing endpoint.

4.3 S3 request accounting and correctness checks

The work metric is observed S3 HEAD + GET + LIST requests during maintenance. Counts continue until processing completes, even if a partition becomes visible earlier. Separate trial buckets isolate maintenance from other methods. CloudWatch request metrics are enabled at least 15 minutes before measurement, and known requests calibrate collection. Two complete quiet minutes follow setup reads; buckets remain idle during a final 10-minute delivery wait. Fixture creation, verification queries, calibration, observer activity, and cleanup are excluded from the maintenance work total.

CloudWatch Sum counters in 60-second periods are corroborated with deduplicated CloudTrail data events matched to bucket, time window, and requester. HEAD includes HeadObject/HeadBucket; GET includes GetObject; LIST includes

ListObjects/ListObjectsV2. Missing measurements remain unknown, not zero.
Both AWS telemetry mechanisms have delivery limitations [7, 8]; agreement and calibration strengthen the observations without making them billing-grade counts.
Of the 192 CloudWatch components, 112 have component datapoints and 80 zeros are inferred from absent component series only when collection succeeded, positive activity was recorded, and known operation counts account for AllRequests.
Those inferred observed zeros also agree with retained CloudTrail events; the inference does not establish complete recording of all service-side activity.
All 192 request values (three components across 64 trials) are available and agree between the graph CSV and the two recorded counting sources. The 448 calibration requests matched. Exact catalog membership, partition values, locations, column definitions, and query rows are verified outside request windows. Verification failures and unattempted conditions are retained as such; they are not replaced silently.
The main run's verifier checked all partition values, locations, and column definitions in memory, retaining summaries, targeted observations, and native query rows. Its full historical catalog snapshots are absent. Complete membership checks for that run therefore rely on frozen verification code and recorded assertions; SELECT rows alone cannot exclude extra empty registrations.
The later graph, activity, and buffer runs save expected and observed full partition sets, locations, and column schemas before assertions. Those saved comparisons and reconstructed query results match in the successful follow-ups.
This repairs future auditability without supplying missing historical snapshots.
For request figures, HEAD/GET/LIST are summed within each repetition before taking the median. Figures 2–5 show medians with observed min–max whiskers. Request figures use logarithmic axes; Figure 3 has linear seconds on Y. Figure 1 plots repetition 1, selected before execution, without averaging binary states. The source CSV retains individual measurements and the report inputs retain outcomes.

4.4 Follow-up evidence and workload protocol

Three separately retained runs extend the primary evaluation. Each has its own source snapshot, settings, request windows, and outcomes; measurements are not merged across implementations or treated as additional identical repetitions.

Evidence set	Execution ID	Scope

Graph validation	run_20260930_113402_14ae	20 comparison trials at N=100/1,000; affected-file counts 1/100/1,000; lifecycle, four crash cases, and eight supporting checks, once each
Activity follow-up	run_20260930_151218_2a18	Four burst schedules × two handlers, four continuous membership-changing cases, and four load cases; one repetition each
Buffer comparison	run_20260930_194451_fcee	Two active-partition pools × two buffer sizes × two repetitions; eight conditions

The graph phase passed, but its first activity child stopped on a producer connection failure: ten conditions passed, one failed, and five were unattempted. The parent remains incomplete_or_failed. The named activity follow-up uses fresh fixtures and retains that failed attempt separately. An earlier fixture-seeding failure also remains recorded. No failed outcome is relabelled as successful.
Burst and continuous fixtures start with 100 registered one-file partitions. The immediate control gives each received event its own eligible intent, including events received in one SQS batch; verification, retries, and delivery completion remain the same. Both methods use the same image and resource size. The measured burst comparison has one repetition and runs RTMSF first, then the immediate control; reversed order is supported by the runner but not measured here.
Continuous membership-changing workloads issue 1 or 3 operations/s for 120 seconds. Three of every four operations upload across 10 or 20 new target partitions; the fourth deletes a distinct seed partition's only file. A separate load family starts with 1,000 registered partitions and keeps their seed files. It appends three unique files, then deletes the oldest appended file, repeatedly. The selected points are (1/s,1,000 targets), (5/s,1,000), (20/s,10), and (20/s,1,000).
At 1/s and 5/s only 90 and 450 distinct upload targets are reached; the configured pool size is not the number of partitions actually exercised at those rates.
The buffer comparison repeats the two 20/s conditions at caps of 100 and 1,000 pending messages, reversing cap order in repetition two. Each condition gets a fresh 0.25-vCPU/512-MiB process with eight table definitions loaded and traffic to one table. Both caps use Q=10 s, T_max=31 s, a 16-MiB payload budget, and a 1,000 s lease that covers the larger cap's conservative renewal budget. The longer lease can delay crash replay. Other reported runs use the 180 s lease; their timings are not pooled with this comparison.

Actual producer completions, scheduling delays, PMCP progress, observed ACKs, approximate SQS counts, and retries are retained. Backlog matches completed file operations to decoded notifications and ACKs, including not-yet-received events; its age assumes synchronized host/worker clocks. Drain starts after the last producer completion and includes telemetry waits and eight seconds of observed queue emptiness. First-partition visibility is independently polled in workloads that create new partitions; it is not a tail-latency measure across all changes.
The activity follow-up passed 16 conditions/6,640 operations, with all 48 request components and 112 calibration calls agreeing. The buffer comparison passed 8 conditions/19,200 operations, with 24 components and 56 calibration calls agreeing. CloudWatch inferred 23 and 16 zero components, respectively, through the same accounting rule as Section 4.3. Graph validation supplied 60 matching components and 140 calibration calls. Later catalog snapshots and exact query rows are retained and audited separately from these request totals.

5. Results

5.1 Membership lifecycle and recovery

The lifecycle starts with an empty catalog and changes one partition: upload its first file, append four, delete one, and delete the remaining four. The resulting file counts are 1, 5, 4, and 0, with desired membership present, present, present, and absent. Each stage waits for processing and verifies the catalog and query rows. Figure 1 shows the interval between a storage change and its observed catalog response. Appending or partly deleting files leaves membership unchanged. The long intervals between stages include test verification; they are not a continuous production workload or per-stage latency benchmark.

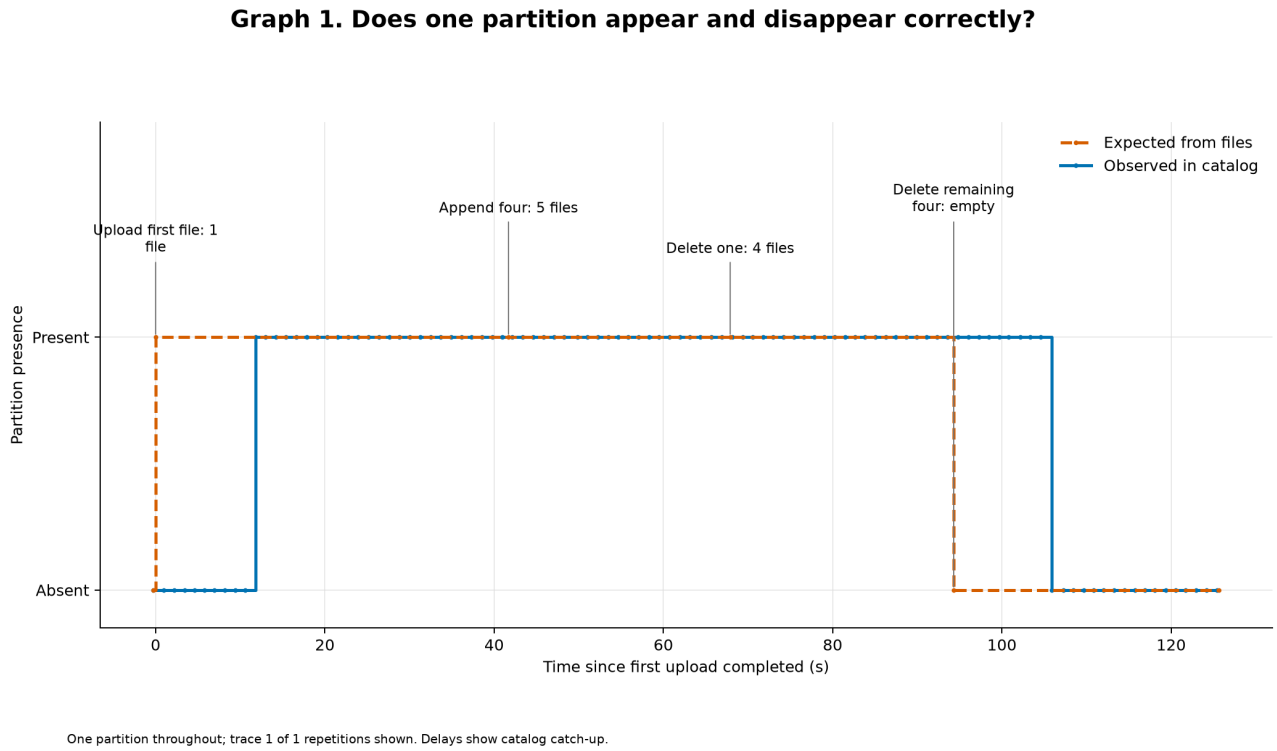


Figure 1. One measured lifecycle, with independent catalog polling. The desired state changes after completed storage operations; observed membership follows after processing. Negative time is the pre-upload observation.

For recovery, a separate experiment launcher wraps the real worker. It targets a partition, operation, and boundary, writes a durable marker outside table data, and calls os._exit(86). The marker suppresses another injected crash after replacement. SQS redelivers unacknowledged input, while the external observer survives. The injection is not part of normal worker startup. Each case completes an empty → registered → empty lifecycle, with one crash at the named boundary.

Table 1. Crash recovery. One attempt per case. State is checked after recovery, before advancing the lifecycle or cleaning up. Verification includes the injected exit, replacement without overlapping workers, same-message replay and acknowledgement, catalog state, and query results.

Crash boundary	Expected after recovery	Observed	Verified crash	Verified replay + ACK	Passed
ADD before PMCP processing	Registered	Registered	1/1	1/1	1/1
ADD after registration, before ACK	Registered	Registered	1/1	1/1	1/1
DROP before PMCP processing	Absent	Absent	1/1	1/1	1/1

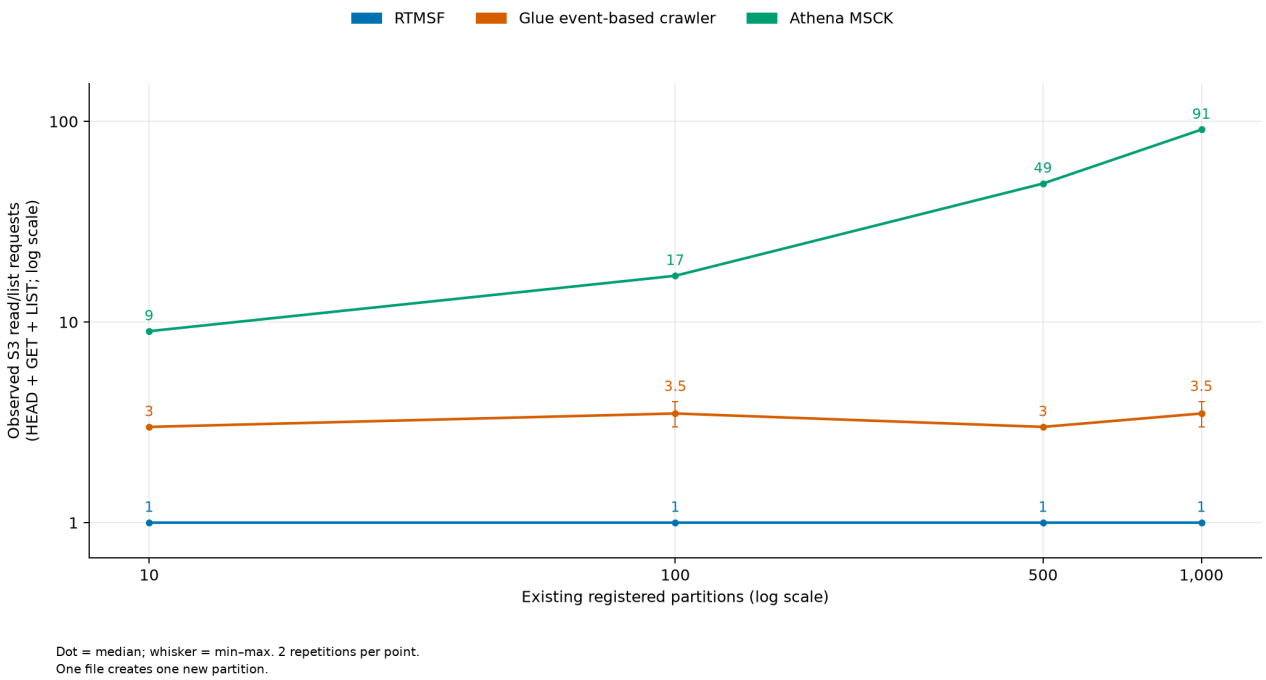
DROP after removal, before ACK	Absent	Absent	1/1	1/1	1/1

All four cases passed with no unknown states. This tests recovery at these boundaries; it is not exhaustive fault injection or validation of concurrent workers. The later graph-validation run also passed the lifecycle and each of these four crash cases once, under its separately retained source and settings.

5.2 Storage work as the table grows

For each starting size, one file is added in a previously absent partition. All methods reach the exact expected N+1 catalog and query result. RTMSF uses one HEAD and no GET/LIST in every measured trial. The event crawler uses 3–4 total requests. Athena uses 9, 17, 49, and 91 requests at 10, 100, 500, and 1,000 starting partitions respectively, identically in both repetitions.

Graph 2. Storage requests as the table grows



Dot = median; whisker = min–max. 2 repetitions per point.
One file creates one new partition.

Figure 2. HEAD + GET + LIST for an isolated new partition in the dt/site/hour layout. Fresh fixtures, two repetitions per method/size; median and min–max.

The event crawler also localizes work at these sizes.

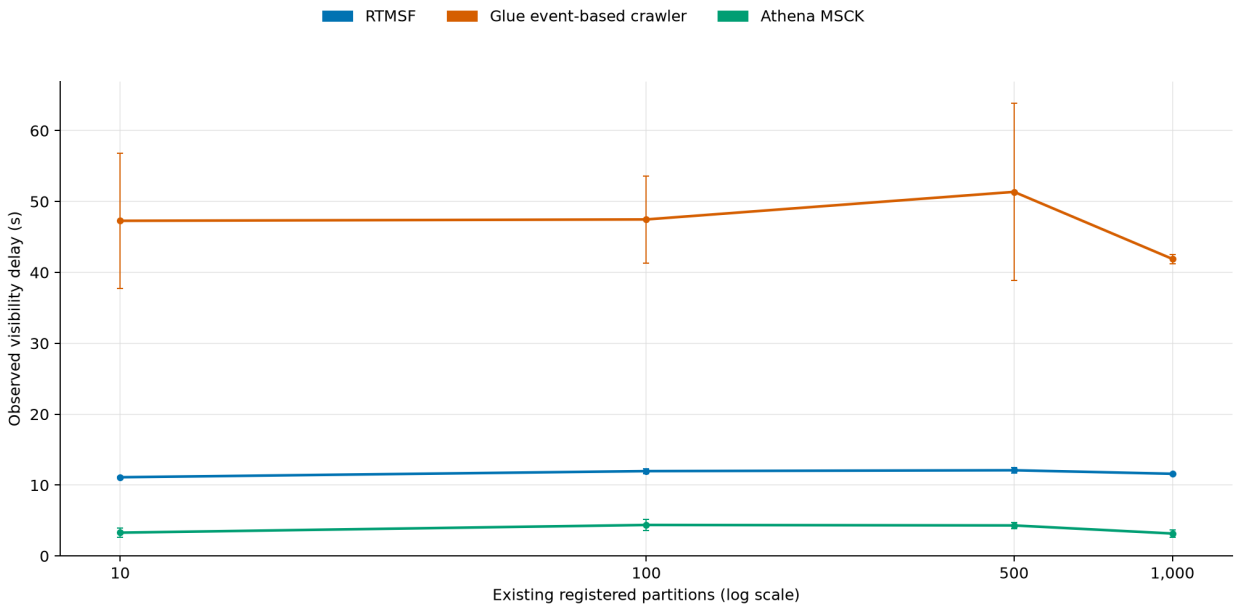
This supports the locality of RTMSF's new-partition check. The small crawler counts also confirm that an event-based crawler should not be described as rescanning every historical partition. Athena's traversal work grows in this

layout, but the count is substantially below one request per partition.

5.3 Response time

Across the eight points per method, RTMSF's observed response times range from 11.008 to 12.464 s, the crawler's from 37.734 to 63.883 s, and Athena's from 2.644 to 5.144 s. Athena is faster than RTMSF at every tested size and repetition under these starts. RTMSF's configured 10-second settling window is included.

Graph 3. Time from each method's trigger to catalog visibility



Dot = median; whisker = min–max. 2 repetitions per point.
Clock starts at upload completion (RTMSF), crawler start request (Glue), or query submission (Athena).

Figure 3. First independently observed catalog presence after each method's trigger (Section 4.2). Two repetitions per method/size; median and min–max.

No scheduled waiting interval is added to Athena or the crawler.

The measured Athena request increase does not yield a monotonic visibility-time increase in this range. Request count alone does not determine elapsed time: service startup, batching, concurrency, and variation can affect the result.

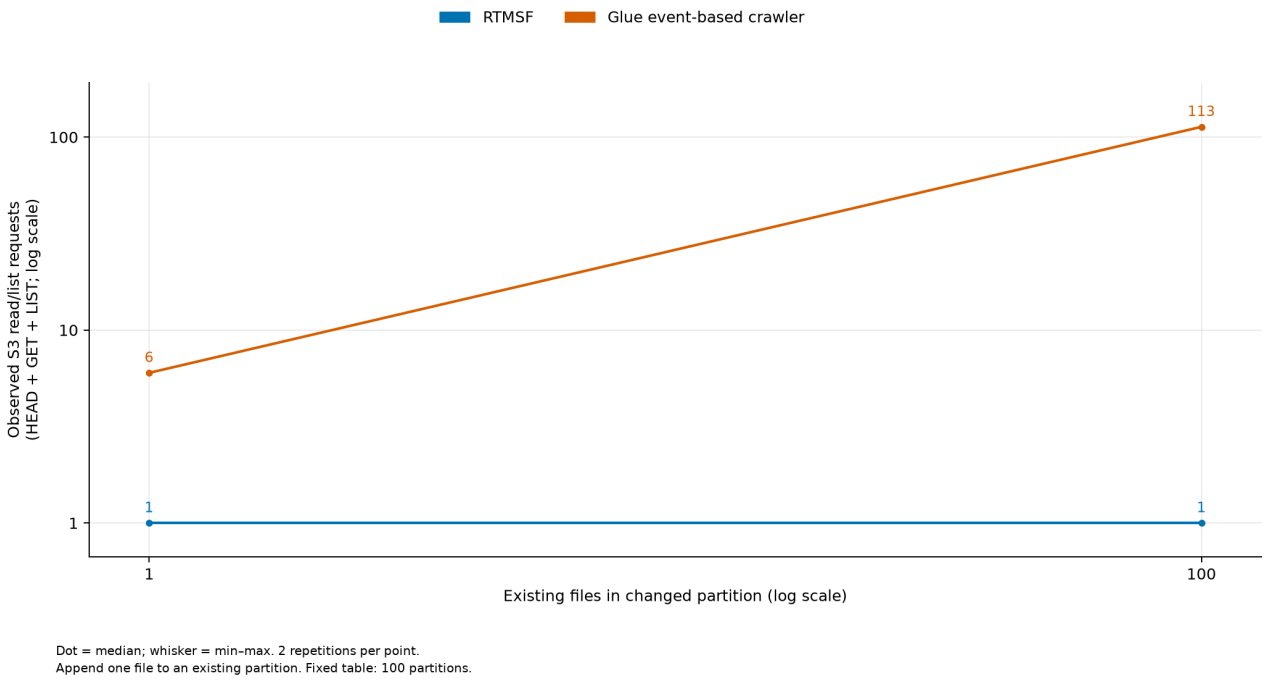
The measurements do not isolate those internal mechanisms. Similarly, the lower crawler median at 1,000 than at 500 partitions, with overlapping observed ranges, is not evidence that larger tables make crawling faster. These results establish the observed latency tradeoff; they do not predict a future size at which RTMSF will overtake Athena or show layout-independent response time.

5.4 Files within the affected partition

Figure 4 appends one file to an already registered partition, with either 1 or 100 existing files and 100 total starting partitions. Membership should remain unchanged, and every trial passes that check and the additional query-row check. This is membership-maintenance work after an append, not the time to expose a new partition: ordinary queries can already discover files within its registered location without another registration.

RTMSF performs one LIST and no HEAD/GET in both conditions and repetitions. The crawler performs 3 HEAD + 2 GET + 1 LIST = 6 requests at one existing file, and 11 HEAD + 101 GET + 1 LIST = 113 requests at 100 existing files. The growing component is principally GET inspection, not one LIST per file.

Graph 4. Storage requests as the affected partition grows



Dot = median; whisker = min–max. 2 repetitions per point.
Append one file to an existing partition. Fixed table: 100 partitions.

Figure 4. Append one file to an already registered partition. X counts files before the append. Two repetitions, 100 starting partitions; median and min–max. Counts cover all observed HEAD/GET/LIST maintenance requests.

These observations expose the value of specializing the operation. RTMSF needs to establish nonemptiness, whereas the configured crawler performs additional file inspection. This does not establish that the crawler's reads are wasted for schema discovery, or that every configuration reads all files. It also does not claim an advantage over deliberately skipping maintenance for an append known to preserve membership. The comparison measures ADD/append maintenance under the recorded configurations. It does not show a comparison between two

implementations of the complete ADD/DROP contract, because the crawler's recorded deletion policy is LOG.

The later graph validation extended the affected-file count to 1,000 with one trial per method/count. Its crawler totals were 6, 103, and 1,038 requests at 1, 100, and 1,000 existing files; GET counts were 2, 101, and 1,001. RTMSF used one LIST at each count. The 100-file crawler total differs from the main run's 113 because HEAD counts vary. This follow-up supports the file-inspection trend without treating the exact total as universal or pooling it into Figure 4.

5.5 Folder-order sensitivity

Figure 5 uses identical tenant/day pairs and file bytes in two layouts: day=2026-09-25/tenant=t000000/ and tenant=t000000/day=2026-09-25/. A fresh fixture starts with N registered partitions; one new tenant is added on the same day. The experiment changes folder order without selecting only an expensive traversal layout for the main comparison.

At 1,000 starting partitions, Athena makes 8 requests in day/tenant order and 2,008 in tenant/day order, in both repetitions. RTMSF makes one HEAD in either order. At 10, 100, and 500 partitions, Athena's corresponding counts are 7, 7, and 7 for day/tenant and 27, 207, and 1,007 for tenant/day.

Graph 5. Does folder order change storage work?

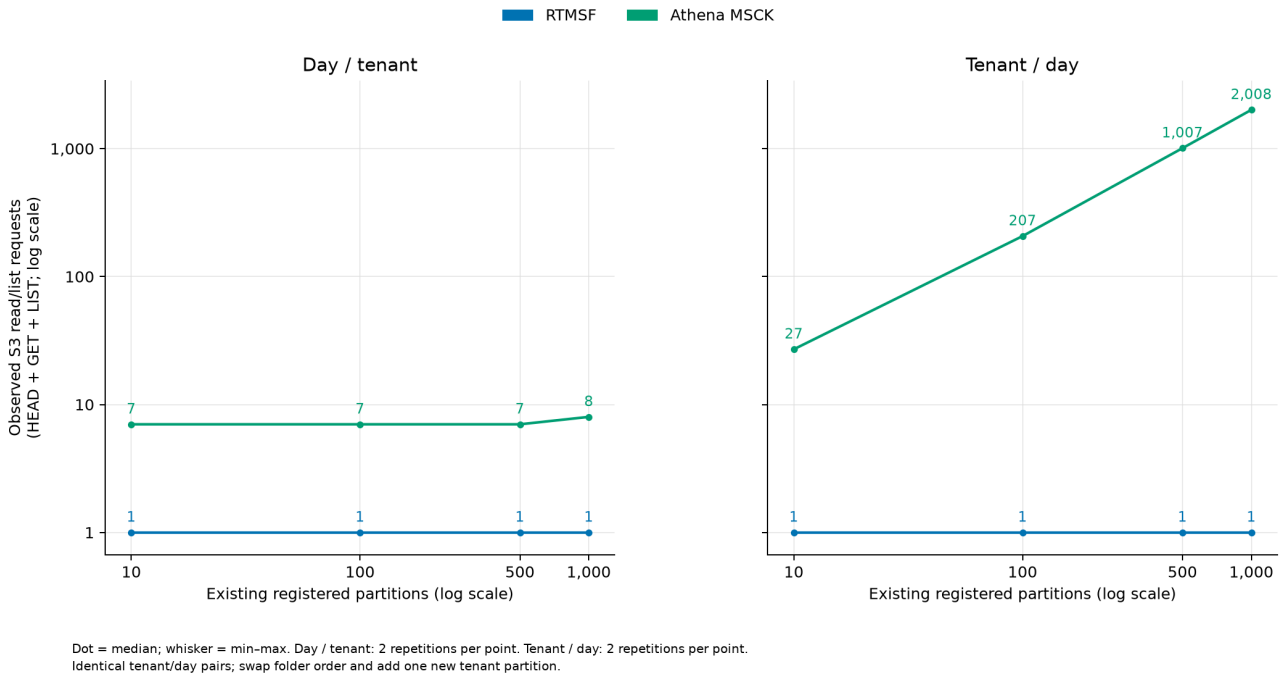


Figure 5. One new tenant/day partition in two folder orders, with shared axes.

Two fresh repetitions per layout/method/size; median and min–max. The same partition count can imply very different discovery work.

Partition count alone therefore does not characterize discovery cost. S3 listing can group prefixes and paginate results [9]; where values occur in the hierarchy changes the traversed namespace. RTMSF's checked object is obtained from the notification, so neither tested order increases its new-partition storage check. This result concerns request work, not latency or a universal lower bound for Hive, Spark, or other repair implementations.

5.6 Supporting checks

The main run's seven additional checks each passed once. The graph-validation follow-up repeated them and added the marker lifecycle; all eight passed once. These test specific behaviors and are not separate comparison curves.

Check	Workload and observed outcome
Short burst	20 distinct file uploads to one partition; 20 observed events, one PMCP attempt, one applied catalog change; 20 rows verified
Long burst	Nine uploads nominally five seconds apart; progress checked at the buffering cap and nine final rows verified
Two tables	30 uploads across three partitions in each of two configured tables; routing and query rows verified
Duplicate and old events	Replayed notifications and old CREATE/DELETE after state changes; final membership verified
Rapid changes	Create/delete and delete/recreate while the worker is stopped; final state and terminal outcomes checked after restart
Retryable failure	Two injected catalog-call failures followed by successful real calls
Retry exhaustion	Initial attempt plus five retries fail; durable dead-letter handling before ACK and subsequent healthy replay verified
Marker lifecycle (follow-up only)	Marker-only registration; add then delete a Parquet file while retaining the marker; delete the marker last and verify removal

These checks support the tested coalescing, routing, and failure behavior. The short burst is not a component-removal experiment and does not imply that every 20 events always collapse to one attempt. Arrival timing, retries, and separate windows can increase that number. Rapid-change success can include the documented stale-CREATE dead-letter outcome; a correct final state is not equivalent to no terminal records.

5.7 Coalescing versus the immediate control

Table 2 reports the activity follow-up's four 20-operation burst schedules.
Both handlers passed their final catalog/query checks. Here each completed PMCP
attempt performs one storage check, so attempts and HEAD+GET+LIST totals coincide.
Table 2. Controlled per-event comparison. One repetition per method/schedule.
Visibility is seconds from first successful upload to the first new partition's
observed registration; it does not measure completion of every change.

Schedule	RTMSF / immediate storage requests	RTMSF / immediate first visibility (s)
20 uploads back-to-back into one new partition	1 / 20	14.404 / 1.388
20 uploads into one new partition, two seconds apart	2 / 20	32.773 / 1.120
One upload into each of 20 new partitions	20 / 20	12.656 / 1.087
Alternate 10 uploads into one new partition with 10 last-file deletions	11 / 20	15.039 / 2.437

Coalescing saves checks when notifications repeatedly concern a partition within
a window; it provides no storage-check reduction for the distributed burst.
The immediate control is faster to first visibility in all four measured cases.
The two-second schedule also forces progress before the quiet interval can finish,
producing two RTMSF attempts. This measures the work/delay tradeoff of enabling
grouping at fixed Q=10/T_max=31. It neither establishes an optimal configuration
nor an advantage over a handler with equivalent grouping and reconciliation rules.
For the motivating dispatcher, an illustrative one-minute discovery budget would
make the quick-burst tradeoff useful: RTMSF used one storage check and exposed
the first partition in 14.404 s, versus 20 checks and 1.388 s for immediate
processing. Both observed times fit that example budget. The budget is illustrative;
the dispatcher itself and a per-partition deadline under ongoing activity are
not evaluated. This connects the observed savings to a possible operating need
without establishing a workload-wide freshness guarantee or total-cost advantage.

5.8 Processing during continuous changes

The four 1/3-operation-per-second membership-changing workloads all progressed
while writes continued, drained, and matched their final catalog/query contents.

At 1/s the 120 operations required 120 attempts for both target pools. At 3/s, 360 operations required 130 attempts for 10 targets and 170 for 20. Respectively, 317 and 321 event observations completed during writes. First-new-partition visibility ranged from 11.068 to 33.384 s across these cases; Q=10 s is not a freshness guarantee. Each case ran once and is not a sustained-capacity benchmark. The load family separates queue behavior from membership transitions by retaining seed files. Table 3 reports the actual achieved rates, unacknowledged operations when writes ended, their oldest age, and observed drain with the default cap of 100 pending deliveries. All cases ended with the expected 1,000 registrations and exact query rows. Correct final state does not mean the worker kept up.

Table 3. Bounded load observations. One approximately 120-second execution per condition. The target pool is 1,000 except in the explicitly marked hot-partition case.

Offered rate / target pool	Achieved operations/s	Unacknowledged at write end	Oldest age (s)	Observed drain (s)
1/s / 1,000	1.007	12	10.941	27.835
5/s / 1,000	4.999	57	11.254	27.663
20/s / 10	19.982	1,400	83.653	180.060
20/s / 1,000	19.986	1,544	89.760	246.364

The two 20/s cases expose backlog beyond the intended settling interval. These measurements describe the supplied serial worker and its configuration, not an algorithmic maximum throughput. At lower rates the smaller exercised partition set limits cross-row capacity comparisons. No producer or worker retries occurred in the successful activity follow-up; one fixture upload recovered before measurement.

5.9 Buffer capacity and catch-up

Increasing the shared pending-message cap improves catch-up in both repeated workloads (Table 4), with a stronger reduction in storage checks for the hot 10-partition pool. All 19,200 operations matched notifications and ACKs; all eight final catalog sets and independently reconstructed query results matched.

Table 4. Buffer comparison at 20 operations/s. Ranges cover two repetitions; both caps use the same 1,000 s lease. Every measured storage request was LIST.

Active partitions	Message cap	LIST requests	Unacknowledged at write end	Observed drain (s)
10	100	244–246	1,500–1,513	194.1–215.6

10	1,000	40	192–354	27.1–32.0
1,000	100	2,318–2,327	1,619–1,642	268.4–274.5
1,000	1,000	2,018–2,081	1,345–1,390	77.6–88.9

Peak process RSS across fresh workers was 60.5–64.2 MiB within a 512-MiB allocation; sampled pending payload remained below 0.81 MiB. A separate local synthetic check filled the 16-MiB raw-payload ceiling, observed receive backpressure, and verified that acknowledgement released capacity. It is a safety check, not AWS performance evidence or a bound on total process memory for arbitrary table configurations.

Two Glue read timeouts recovered through PMCP retries in the second larger-buffer, 1,000-target trial: 2,020 attempts produced 2,018 LIST calls. Two observation reads also retried; there were no producer retries or DLQ messages. The larger buffer did not eliminate backlog in the distributed workload or establish sustained 20/s service. Queue retention absorbs finite bursts; permanently exceeding processing capacity increases lag even when worker memory is bounded.

The cap counts retained notifications, rather than partition groups. When 100 notifications fill it, the worker stops receiving even if those notifications concern only ten partitions. Input waiting in SQS cannot join a grouping window until received. After buffered work completes and its deliveries are acknowledged, later input can open new windows and cause repeated checks. The larger cap admits more related input before completion: the hot-partition workload averages 60 notifications per storage check, versus approximately ten with the smaller cap. Across 1,000 active partitions, less activity can be combined and separate checks remain necessary. The observed lesson is that intake capacity and partition locality jointly determine grouping effectiveness at fixed Q and T_max.

Pending work at the end of writing also needs an age measure. With the larger cap, the oldest unacknowledged operation was 9.886–17.857 s old for ten targets, compatible with intentional grouping, versus 73.297–77.597 s for 1,000 targets. The latter exceeds the initial grouping interval and indicates accumulated lag.

In the first larger-buffer, 1,000-target trial, task-scoped telemetry records 2,081 catalog reads, 2,081 LIST calls, and 2,400 acknowledgement requests. Their SDK durations sum to 74.632, 41.609, and 13.315 s respectively, approximately 130 s of serial remote calls for a workload produced in 120 s. Receiving and

other worker activity add time. This accounting explains a limit in the supplied serial implementation; increasing buffer space alone does not increase its execution concurrency. It is not a maximum-capacity estimate or a predicted speedup for a parallel worker.

6. Discussion and limitations

The evidence favors selecting maintenance mechanisms by required semantics and namespace shape. For an isolated addition with an immediately issued command, Athena is a fast option in the measured range. For workflows requiring schema discovery, a crawler supplies capabilities outside RTMSF's purpose. For known-schema membership with independent writers and deletion activity, RTMSF combines local checks and retained-work recovery while keeping the core independent of AWS APIs.

Reduced S3 requests are not measured lower total cloud cost. SQS traffic, catalog calls, worker uptime, logging, crawler service charges, and operational maintenance also matter. No throughput or cost crossover is established. Nor does a multiple-table deployment force an event-aware MSCK scheduler to repair every table; a scheduler can target affected tables. The two-table check demonstrates routing correctness, not large-fleet efficiency.

The main figures have four small partition counts and two comparison repetitions. Separate graph validation extends the affected-file count to 1,000 and repeats lifecycle/fault checks once. The immediate-handler comparison has one repetition; the focused buffer comparison has two. All activity streams last about two minutes. These observations do not establish stable latency distributions, maximum sustainable throughput, 100K/500K scaling, long-run reliability, or production adoption. A Q sweep is not performed: the fixed 10-second setting demonstrates a tradeoff at that configuration, not an optimum. AWS services and engines can change; broader runs may revise conclusions as well as numbers.

There is identifiable implementation overhead beyond the core algorithms. In both 1,000-target, 1,000-buffer trials, worker telemetry records 2,400 SQS DeleteMessageBatch calls for 2,400 notifications: the worker currently passes one completed delivery at a time. SQS accepts up to ten per request and reports per-entry failures [13]. Grouping completed acknowledgements could reduce queue calls; another cloud's adapter could instead issue individual acknowledgements.

The supplied worker also processes partition checks serially. ACK batching and bounded concurrency across different partitions are deferred engineering improvements, not evaluated speedups or new WOERA/PMCP contributions. The paper reports the implementation as measured and does not assume those changes would eliminate backlog. Buffer tuning alone does not supply unlimited service capacity. Notifications can be delayed or permanently omitted. Replay only repairs work that remains admitted and retained; an independent repair procedure may be an operational backstop. A single worker does not fence arbitrary writers, and an emptiness observation cannot make catalog removal atomic with new writes. Multi-worker fencing, table/schema changes, and long-running overload are unvalidated. Stable routing by table/partition is necessary for exclusive ownership; sharing this Standard-SQS queue across workers alone does not establish it. Marker-only prefixes deliberately count as nonempty (Section 3.1). These are boundaries of the system, not guarantees supplied by the crash tests.
Finally, only the AWS adapters are evaluated. Implementing another provider requires checking notification, retention, storage, and catalog semantics as well as translating APIs. Provider-neutral interfaces alone do not demonstrate cross-cloud conformance. This evaluation does not claim that established recovery or buffering mechanisms are individually novel, or that the prior-work search is exhaustive.

7. Conclusion

RTMSF integrates event routing, partition-level grouping, bounded storage checks, and retained-message replay into a provider-neutral membership-maintenance system for existing Hive-style tables. Writers continue using ordinary object-store operations. In the measured AWS workloads, bounded storage checks avoid broader namespace discovery and per-file inspection for this task. The immediate-handler comparison shows that coalescing saves checks when partition activity repeats within a window, while exchanging earlier visibility for less work. The distributed burst has no such reduction. Lifecycle, crash, marker, and active-load checks exercise the stated processing and recovery contract. Buffer validation demonstrates safe catch-up in the tested finite workloads, with observed overload rather than a claim of unlimited throughput. Athena's and the immediate handler's lower

response times reinforce the distinction between storage work and latency.

The evidence supports workload-dependent request reduction under explicit notification, retention, ownership, and concurrency assumptions.

References

- Amazon Web Services. MSCK REPAIR TABLE — Amazon Athena. https://docs.aws.amazon.com/athena/latest/ug/msck-repair-table.html
- Amazon Web Services. Accelerating crawls using Amazon S3 event notifications. https://docs.aws.amazon.com/glue/latest/dg/crawler-s3-event-notifications.html
- Michael Armbrust et al. Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores. PVLDB 13(12), 2020. https://www.vldb.org/pvldb/vol13/p3411-armbrust.pdf
- Sagar Sumit, Prashant Wason, and Sivabalan Narayanan. Metadata-as-Data in Apache Hudi: A Multi-Modal Index Substrate for Lakehouse Tables. Proceedings of the 1st International Workshop on Data FORMATS for Modern Architectures and Workloads, pp. 1–5, 2026. https://doi.org/10.1145/3802514.3809177
- Keonwoo Oh, Pooja Nilangekar, and Amol Deshpande. TreeCat: Standalone Catalog Engine for Large Data Systems. PVLDB 18(11), pp. 4323–4336, 2025. https://www.vldb.org/pvldb/vol18/p4323-oh.pdf
- Tyler Akidau et al. MillWheel: Fault-Tolerant Stream Processing at Internet Scale. PVLDB 6(11), 2013. https://www.vldb.org/pvldb/vol6/p1033-akidau.pdf
- Amazon Web Services. Monitoring metrics with Amazon CloudWatch — Amazon S3. https://docs.aws.amazon.com/AmazonS3/latest/userguide/cloudwatch-monitoring.html
- Amazon Web Services. How AWS CloudTrail works. https://docs.aws.amazon.com/awscloudtrail/latest/userguide/how-cloudtrail-works.html
- Amazon Web Services. ListObjectsV2 — Amazon S3 API Reference. https://docs.aws.amazon.com/AmazonS3/latest/API/API_ListObjectsV2.html
- Snowflake. Refresh external tables automatically for Amazon S3. https://docs.snowflake.com/en/user-guide/tables-external-s3
- Snowflake. Introduction to external tables. https://docs.snowflake.com/en/user-guide/tables-external-intro

• Amazon Web Services. Partition projection with Amazon Athena.
https://docs.aws.amazon.com/athena/latest/ug/partition-projection.html
• Amazon Web Services. DeleteMessageBatch — Amazon SQS API Reference.
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/APIReference/API_DeleteMessageBatch.html

Artifact and reproduction

The repository contains the manuscript, five figures, plotting CSVs, recorded report inputs and a source-provenance index. The companion full evidence archive adds per-execution source archives, immutable image/source identities, scoped request and calibration records, operation/ACK telemetry, query checks and later catalog snapshots. Failures and historical missing measurements remain visible. The archive is prepared for release; this manuscript does not assert an assigned arXiv identifier or a public data DOI.

• Paper evidence and regeneration commands.
• Primary graph data and report inputs.
• Source identities, including the primary source-manifest digest b8623cf17cb6c9b2647717dc467cb05c232391cfa96d4ae5346ff724f8116efc. This identifies the sorted source-file hash manifest, not compressed archive bytes.
• Activity inputs, buffer inputs, and buffer call accounting.
• Experiment specification, system behavior, and AWS runbook.

From a clean checkout, install the locked dependencies once and regenerate offline:

```
uv sync --frozen --extra dev
uv run --frozen --extra dev pytest -q
uv run --frozen --extra experiment python -m experiments.figures \
  --csv evidence/paper/figures/graph_data.csv \
  --results evidence/paper/figures/report_inputs.json \
  --output-dir runs/regenerated-figures
uv run --frozen --extra experiment python -m experiments.activity_reports \
  --results evidence/paper/activity/report_inputs.json \
```

```
--output-dir runs/regenerated-activity
```

Use the buffer or failed-activity input to regenerate those reports separately.
No AWS account is needed for reporting or local tests. The compact repository inputs omit verbose audit envelopes; the full archive retains the supporting records. Account identifiers use a consistent placeholder; credential identities, receipt handles and source IPs are omitted. Source archives retain their original bytes. SHA256SUMS verifies the exported artifact; original artifact lists refer to original private bytes and are labelled accordingly.
The runbook gives exact size/repetition options for each measured matrix and AWS setup, image validation and cleanup. Every execution uses fresh fixtures, retains its own source version, CSV, outcomes and failures, and generates reports without manual assembly. Exact historical formatting uses that execution's archived source. New cloud runs seek comparable conclusions, not identical timing.
Neither the diagnostic Athena-only probes nor older 10K runs supply the five figures above. Larger-scale evidence and additional cloud implementations are separate extensions; their outcomes cannot be assumed in advance.